

# Lab 2: Primary–Backup

**MSC5703/MCS4993: Intro to Distributed Computing Fall 2025**  
**Duration: Two Weeks Due Date: [Nov 5, 2025]**

## 1. Introduction & Objectives

In this lab, you will build a simple, fault-tolerant service using the Primary–Backup replication scheme. This is a common and fundamental technique for building systems that can survive the failure of a single server.

Your goal is to implement a fault-tolerant Key-Value (KV) store. Clients should be able to `Put` a value associated with a key and `Get` the value for a key. The service must continue to operate correctly even if one of the servers (either the primary or the backup) crashes.

To keep this lab “simple,” we will make a few key assumptions:

- We will not handle network partitions (i.e., you can assume the network is reliable).
- We will not handle a “split-brain” scenario.
- We will use a central `View Service` to act as the single source of truth for which server is the primary and which is the backup. This `View Service` is a single point of failure (SPOF), which is acceptable for this lab.

### Learning Goals:

- Understand the design and trade-offs of the Primary–Backup replication model.
- Implement a fault-tolerant service that handles server failures.
- Gain practical experience with RPC (Remote Procedure Call).
- Manage distributed state and handle state transfer.

## 2. System Architecture

Your system will consist of three components:

**The `View Service`:** A single, central server that maintains the “view” of the system.

- It decides who is the `Primary` and who is the `Backup`.
- It detects server failures via heartbeating (`Ping`s).
- It manages the failover process by promoting the backup.

**The KV Servers ( Primary & Backup ):** These are the servers that store the actual data (the key-value map).

- Only one server is the `Primary` at any given time.
- Only the `Primary` can accept `Put / Get` requests from clients.
- The `Backup`'s job is to be a perfect replica of the `Primary`, ready to take over.
- Both servers will periodically `Ping` the `View Service` to announce they are alive.

**The Client:** A simple library (that you will also write) that talks to the service.

- The client first asks the `View Service`, “Who is the primary?”
- It then sends all `Get` and `Put` requests to that primary.
- If its requests fail (e.g., connection refused), it asks the `View Service` again for the new primary.

### 3. Implementation Details

You are free to use any language, but Go, Python, or Java are recommended. You must use an RPC framework for all communication (e.g., Go’s `net/rpc`, gRPC, or Python’s `xmlrpc`). Do not use simple sockets.

#### Part A: The View Service

The `View Service` is the “brains” of the operation. It doesn’t store any key-value data, only the system’s current configuration.

It must expose (at least) these RPCs:

- `Ping (server_name):`
  - Called by KV servers (e.g., “server1”, “server2”) every 0.5 seconds.
  - The `View Service` records the time of the last successful ping from `server_name`.
  - It returns the current `View` object (see below).
- `GetView ():`
  - Called by clients to find the current primary.
  - It returns the current `View` object.

The `View` object should contain:

- `ViewNumber` : An integer that increments every time the view changes.
- `Primary` : The string name/address of the primary server.
- `Backup` : The string name/address of the backup server (can be empty).

**Internal Logic:** The `View Service` must run a “ticker” function (e.g., once every 0.5 seconds) that:

- Checks the last ping time for all servers.
- If the `Primary` hasn’t pinged in (e.g.) 1.5 seconds, it is declared dead.
- If the `Backup` hasn’t pinged in 1.5 seconds, it is declared dead.
- Handles promotions:
  - **Primary failure:** If `Primary` is dead and `Backup` is alive, promote the `Backup` to `Primary`. The view now has a `Primary` but no `Backup`.
  - **Backup failure:** If `Backup` is dead, the view now has a `Primary` but no `Backup`.
  - **New server:** If a new, “idle” server starts pinging and there is no `Primary`, promote it to `Primary`.
  - **New server (as backup):** If a new, “idle” server starts pinging, there is a `Primary`, but no `Backup`, assign the new server as the `Backup`.

This logic ensures the system “heals” itself by assigning roles to available servers.

## Part B: The KV Server (`Primary` & `Backup`)

This is the core of your lab. You will write a single server program that can act as either a `Primary` or a `Backup`, depending on what the `View Service` tells it.

### Server State:

- Its own name/address (e.g., “server1”).
- The current `View` it got from the `View Service`.
- The key-value data, stored in a simple hash map (e.g., `map[string]string`).

**Main Loop:** In a background thread, the KV server must:

- Call `Ping()` on the `View Service` every 0.5 seconds.
- Receive the new `View` from the `Ping()` response.
- Compare the new `View` with its current `View`.
- If the `View` has changed, update its internal state and role.

### Role-Based Logic:

*If I am the PRIMARY:*

- Accept `Get` (key) RPCs: Look up the key in the local map and return the value.

- Accept Put (key, value) RPCs:
  - Check if there is a Backup in the current view.
  - If yes, first forward the Put request to the Backup (e.g., via an internal ForwardUpdate (key, value) RPC).
  - Wait for the Backup to acknowledge (return OK) that it has stored the data.
  - Only after the Backup acks, update the Primary's own local map.
  - Reply OK to the client.
  - If no Backup exists, just update the local map and reply OK to the client.

This is synchronous replication and ensures the Backup is never behind the Primary.

*If I am the BACKUP:*

- Reject all Get (key) and Put (key, value) RPCs from clients (e.g., return an “I am not primary” error).
- Accept the internal ForwardUpdate (key, value) RPC from the Primary.
- On ForwardUpdate, update the local map and return OK to the Primary.

### Part C: State Transfer

There is one critical challenge: What happens when a new Backup joins the system?

- The View Service will assign it as the Backup.
- The Primary will see this new Backup in the View it gets from its Ping().
- This new Backup is empty! It doesn't have the Primary's data.

You must implement state transfer:

- When the Primary sees a new server has become the Backup (e.g., View.Backup was empty, and now it's “server2”), it must transfer its entire state to the new Backup.
- The Primary should send its entire key-value map to the Backup (e.g., via a new internal RPC like SyncState (map[string]string)).
- The Backup receives this map, overwrites its local map, and acks.
- Crucially: While the state transfer is happening, the Primary must queue (but not process) any new client Put requests to avoid inconsistency. It can process them after the transfer is complete.

## 4. The Client

Your client library should expose two functions:

- `Get` (`key`)
- `Put` (`key, value`)

Internally, the client must:

- Keep track of the last known `Primary`.
- On a `Put` or `Get`, try to send the RPC to that `Primary`.
- If the RPC fails (e.g., server is dead, or it replies “I am not primary”), the client must:
  - Call `GetView()` on the `View Service` to get the new `Primary`.
  - Retry the `Put / Get` request with the new `Primary`.
  - Loop until it succeeds.

## 5. Testing

You are responsible for testing your own code. We will run a similar test harness to grade your lab. A good test script would:

1. Start the `View Service`.
2. Start KV server “S1”.
3. Check the view (S1 should be `Primary`).
4. Client: `Put (“a”, “1”)`, `Get (“a”)` (should return “1”).
5. Start KV server “S2”.
6. Check the view (S1 `Primary`, S2 `Backup`).
7. Client: `Put (“b”, “2”)`.
8. **Test Primary Failure:** Kill S1.
9. Wait 2–3 seconds.
10. Check the view (S2 should be `Primary`, no `Backup`).
11. Client: `Get (“a”)` (should return “1”), `Get (“b”)` (should return “2”).
12. **Test State Transfer:** Start KV server “S3”.
13. Check the view (S2 `Primary`, S3 `Backup`).

14. Wait for state transfer to complete.
15. **Test New Backup:** Kill S2.
16. Check the view (S3 Primary, no Backup).
17. Client: Get ("a") (should return "1"), Get ("b") (should return "2").

If your client gets the correct data after all these failures, your implementation is correct!

## 6. Deliverables

**Source Code:** A zip file containing your complete, commented source code for the View Service, the KV Server, and the Client library.

**Report:** A report includes:

- Your name and student ID.
- A description of your design, particularly any challenges you faced (e.g., how you handled state transfer).
- Results and analysis.
- Any known bugs or limitations.
- Instructions on how to compile and run your code.

## 7. Grading Policy (100 points total)

- **View Service correctness (20 pts):** Accurate ping tracking; timely detection of dead servers; correct promotion/demotion; monotonic ViewNumber updates.
- **KV Primary/Backup logic (30 pts):** Primary accepts client Put / Get ; Backup rejects client calls; synchronous replication on Put with waiting for Backup ACK before commit; correct local updates and client replies.
- **State Transfer (25 pts):** Full snapshot via SyncState ; Backup state overwrite; Primary queues client Put requests during transfer and drains afterward; no inconsistencies.
- **Client & Report (25 pts):** Client retries with GetView on role errors/failures until success; clear Report with build/run instructions and concise design notes.